

## ✓ Dependencies

```
import numpy as np
import tensorflow as tf
import tensorflow_datasets as tfds
import matplotlib.pyplot as plt
gpus = tf.config.experimental.list_physical_devices('GPU')
for gpu in gpus:
    tf.config.experimental.set_memory_growth(gpu, True)
```

## ✓ Loading the data

```
data = tfds.load('fashion_mnist', split='train')
```

⇨ Downloading and preparing dataset 29.45 MiB (download: 29.45 MiB, generated: 36.42 MiB, D1 Completed...: 0 url [00:00, ? url/s]  
 D1 Size...: 0 MiB [00:00, ? MiB/s]  
 Extraction completed...: 0 file [00:00, ? file/s]  
 Generating splits...: 0%| | 0/2 [00:00<?, ? splits/s]  
 Generating train examples...: 0%| | 0/60000 [00:00<?, ? examples/s]  
 Shuffling fashion\_mnist-train.tfrecord...: 0%| | 0/60000 [00:00<?, ? examples/s]  
 Generating test examples...: 0%| | 0/10000 [00:00<?, ? examples/s]  
 Shuffling fashion\_mnist-test.tfrecord...: 0%| | 0/10000 [00:00<?, ?

```
data.as_numpy_iterator().next().keys()
```

⇨ dict\_keys(['image', 'label'])

## ✓ Data Visualization

```
dataiterator = data.as_numpy_iterator()
```

```
dataiterator.next()
```

⇨

```

[ 1],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[153],
[ 78],
[106],
[ 37],
[ 0],
[ 0],
[ 0]],


[[ 0],
[ 0],
[ 0],
[ 18],
[ 61],
[ 41],
[103],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[ 0],
[106],
[ 47],
[ 69],
[ 23],
[ 0],
[ 0],
[ 0]]], dtype=uint8),
'label': 2}

```

```

fig, ax =plt.subplots(ncols= 4, figsize=(20, 20))
for idx in range(4):
    sample = dataiterator.next()
    ax[idx].imshow(np.squeeze(sample['image']))
    ax[idx].title.set_text(sample['label'])

```



## ✓ preprocessing

```
def scale_images(data):
    image = data['image']
    return image / 255
```

### Preparing data for tensorflow

```
data = data.map(scale_images)
data = data.cache()
data = data.shuffle(60000)
data = data.batch(128)
data = data.prefetch(64)
```

```
print(data)
```

```
<PrefetchDataset element_spec=TensorSpec(shape=(None, 28, 28, 1), dtype=tf.float32, name=
```

```
data.as_numpy_iterator().next().shape
```

```
(128, 28, 28, 1)
```

## ✓ Neural Network

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Dense, Flatten, Reshape, LeakyReLU, Dropout, UpS
```

## ✓ Generator

```
def build_generator():
    model = Sequential()

    model.add(Dense(7*7*128, input_dim= 128))
    model.add(LeakyReLU(0.2))
    model.add(Reshape((7,7,128)))

    model.add(UpSampling2D())
    model.add(Conv2D(128, 5, padding= 'same'))
    model.add(LeakyReLU(0.2))

    model.add(UpSampling2D())
    model.add(Conv2D(128, 5, padding= 'same'))
    model.add(LeakyReLU(0.2))

    model.add(Conv2D(128, 4, padding= 'same'))
    model.add(LeakyReLU(0.2))

    model.add(Conv2D(128, 4, padding= 'same'))
    model.add(LeakyReLU(0.2))

    model.add(Conv2D(1, 4, padding= 'same', activation= 'sigmoid'))

    return model
```

```
generator = build_generator()
```

```
generator.summary()
```



Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 6272)	809088
leaky_re_lu (LeakyReLU)	(None, 6272)	0
reshape (Reshape)	(None, 7, 7, 128)	0
up_sampling2d (UpSampling2D)	(None, 14, 14, 128)	0
conv2d (Conv2D)	(None, 14, 14, 128)	409728

leaky_re_lu_1 (LeakyReLU)	(None, 14, 14, 128)	0
up_sampling2d_1 (UpSampling 2D)	(None, 28, 28, 128)	0
conv2d_1 (Conv2D)	(None, 28, 28, 128)	409728
leaky_re_lu_2 (LeakyReLU)	(None, 28, 28, 128)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	262272
leaky_re_lu_3 (LeakyReLU)	(None, 28, 28, 128)	0
conv2d_3 (Conv2D)	(None, 28, 28, 128)	262272
leaky_re_lu_4 (LeakyReLU)	(None, 28, 28, 128)	0
conv2d_4 (Conv2D)	(None, 28, 28, 1)	2049


```

=====
Total params: 2,155,137
Trainable params: 2,155,137
Non-trainable params: 0

```

---

```
img = generator.predict(np.random.randn(4, 128, 1))
```


 1/1 [=====] - 7s 7s/step

```
fig, ax = plt.subplots(ncols= 4, figsize=(20, 20))
```

```

for idx, img in enumerate(img):
    ax[idx].imshow(np.squeeze(img))
    ax[idx].title.set_text(idx)

```



## ▼ discriminator

```
def build_discriminator():

    model = Sequential()

    model.add(Conv2D(32, 5, input_shape = (28, 28, 1)))
    model.add(LeakyReLU(0.2))
    model.add(Dropout(0.4))

    model.add(Conv2D(64, 5))
    model.add(LeakyReLU(0.2))
    model.add(Dropout(0.4))

    model.add(Conv2D(128, 5))
    model.add(LeakyReLU(0.2))
    model.add(Dropout(0.4))

    model.add(Conv2D(256, 5))
    model.add(LeakyReLU(0.2))
    model.add(Dropout(0.4))

    model.add(Flatten())
    model.add(Dropout(0.4))
    model.add(Dense(1, activation= 'sigmoid'))

    return model
```

```
discriminator = build_discriminator()
```

```
discriminator.summary()
```

➡ Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_5 (Conv2D)	(None, 24, 24, 32)	832
leaky_re_lu_5 (LeakyReLU)	(None, 24, 24, 32)	0
dropout (Dropout)	(None, 24, 24, 32)	0
conv2d_6 (Conv2D)	(None, 20, 20, 64)	51264
leaky_re_lu_6 (LeakyReLU)	(None, 20, 20, 64)	0
dropout_1 (Dropout)	(None, 20, 20, 64)	0
conv2d_7 (Conv2D)	(None, 16, 16, 128)	204928
leaky_re_lu_7 (LeakyReLU)	(None, 16, 16, 128)	0
dropout_2 (Dropout)	(None, 16, 16, 128)	0
conv2d_8 (Conv2D)	(None, 12, 12, 256)	819456
leaky_re_lu_8 (LeakyReLU)	(None, 12, 12, 256)	0
dropout_3 (Dropout)	(None, 12, 12, 256)	0
flatten (Flatten)	(None, 36864)	0
dropout_4 (Dropout)	(None, 36864)	0
dense_1 (Dense)	(None, 1)	36865
=====		
Total params: 1,113,345		
Trainable params: 1,113,345		
Non-trainable params: 0		
=====		

```
img.shape
```

➡ (28, 28, 1)

```
img = generator.predict(np.random.randn(4, 128, 1))
```

➡ 1/1 [=====] - 0s 20ms/step

```
discriminator.predict(img)
```

```
⇒ 1/1 [=====] - 0s 158ms/step
array([[0.49692342],
       [0.49718642],
       [0.49706194],
       [0.49711314]], dtype=float32)
```

## ✓ Custom Training Loop

### Losses and optimizers

```
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import BinaryCrossentropy
```

```
g_opt = Adam(learning_rate= 0.0001)
d_opt = Adam(learning_rate= 0.00001)
g_loss = BinaryCrossentropy()
d_loss = BinaryCrossentropy()
```

### subclasses Model

```
from tensorflow.keras.models import Model
class imageGAN(Model):
    def __init__(self, generator, discriminator, *args, **kwargs):
        super().__init__(*args, **kwargs)

        self.generator = generator
        self.discriminator = discriminator

    def compile(self, g_opt, d_opt, g_loss, d_loss, *args, **kwargs):
        super().compile(*args, **kwargs)

        self.g_opt = g_opt
        self.d_opt = d_opt
        self.g_loss = g_loss
        self.d_loss = d_loss

    def train_step(self, batch):
        real_images = batch
        fake_images = self.generator(tf.random.normal((128, 128, 1)), training= False)

        #Training Discriminator
        with tf.GradientTape() as d_tape:

            yhat_real = self.discriminator(real_images, training= True)
```



```

yhat_fake = self.discriminator(fake_images, training= True)
yhat_realfake = tf.concat([yhat_real, yhat_fake], axis=0)

y_realfake = tf.concat([tf.zeros_like(yhat_real), tf.ones_like(yhat_fake)], axis=0)

noise_real = 0.15*tf.random.uniform(tf.shape(yhat_real))
noise_fake = -0.15*tf.random.uniform(tf.shape(yhat_fake))
y_realfake += tf.concat([noise_real, noise_fake], axis=0)

total_d_loss = self.d_loss(y_realfake, yhat_realfake)

d_grad = d_tape.gradient(total_d_loss, self.discriminator.trainable_variables)
self.d_opt.apply_gradients(zip(d_grad, self.discriminator.trainable_variables))

#Training Generator
with tf.GradientTape() as g_tape:

    gen_images = self.generator(tf.random.normal((128, 128, 1)), training= True)

    predicted_labels = self.discriminator(gen_images, training= False)

    total_g_loss = self.g_loss(tf.zeros_like(predicted_labels), predicted_labels)

    ggrad = g_tape.gradient(total_g_loss, self.generator.trainable_variables)
    self.g_opt.apply_gradients(zip(ggrad, self.generator.trainable_variables))

    return {"d_loss ": total_d_loss, "g_loss ": total_g_loss}

```

instantiating the imageGAN subclassed model

```
imggan = imageGAN(generator, discriminator)
```

compiling the model

```
imggan.compile(g_opt, d_opt, g_loss, d_loss)
```

## ✓ Building Callback

```

import os
from tensorflow.keras.preprocessing.image import array_to_img
from tensorflow.keras.callbacks import Callback

```

```

class ModelMonitor(Callback):
    def __init__(self, num_img=3, latent_dim= 128):
        self.num_img = num_img

```

```

self.latent_dim = latent_dim

def on_epoch_end(self, epoch, logs= None):
    random_latent_vectors = tf.random.uniform((self.num_img, self.latent_dim, 1))
    generated_images = self.model.generator(random_latent_vectors)
    generated_images *= 255
    generated_images.numpy()
    for i in range(self.num_img):
        img = array_to_img(generated_images[i])
        img.save(os.path.join(f'generated_img_{epoch}_{i}.png'))

```

## ✓ Training

```

%%time
hist = imggan.fit(data, epochs= 400, callbacks=[ModelMonitor()])

```

```

⇒ Epoch 1/400
469/469 [=====] - 73s 146ms/step - d_loss : 0.6111 - g_loss : 
Epoch 2/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.3994 - g_loss : 
Epoch 3/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.3042 - g_loss : 
Epoch 4/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.2877 - g_loss : 
Epoch 5/400
469/469 [=====] - 71s 152ms/step - d_loss : 0.2801 - g_loss : 
Epoch 6/400
469/469 [=====] - 71s 152ms/step - d_loss : 0.2758 - g_loss : 
Epoch 7/400
469/469 [=====] - 71s 152ms/step - d_loss : 0.2754 - g_loss : 
Epoch 8/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.3796 - g_loss : 
Epoch 9/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6565 - g_loss : 
Epoch 10/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.4821 - g_loss : 
Epoch 11/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.3381 - g_loss : 
Epoch 12/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.3236 - g_loss : 
Epoch 13/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.5460 - g_loss : 
Epoch 14/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6809 - g_loss : 
Epoch 15/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.6546 - g_loss : 
Epoch 16/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6825 - g_loss : 
Epoch 17/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.6719 - g_loss : 
Epoch 18/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6844 - g_loss : 

```


```
Epoch 19/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6800 - g_loss :
Epoch 20/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6764 - g_loss :
Epoch 21/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6821 - g_loss :
Epoch 22/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.6816 - g_loss :
Epoch 23/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.6913 - g_loss :
Epoch 24/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.6966 - g_loss :
Epoch 25/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.6856 - g_loss :
Epoch 26/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.6720 - g_loss :
Epoch 27/400
469/469 [=====] - 72s 154ms/step - d_loss : 0.6750 - g_loss :
Epoch 28/400
469/469 [=====] - 72s 153ms/step - d_loss : 0.4737 - g_loss :
Epoch 29/400
```

## ✓ Performance Review

```
plt.suptitle('Loss')
plt.plot(hist.history['d_loss'], label='d_loss')
plt.plot(hist.history['g_loss'], label='g_loss')
plt.legend()
plt.show()
```



## Loss




## ✓ Generating images using trained model

```
imgs = generator.predict(tf.random.normal((16, 128, 1)))
```



```
1/1 [=====] - 0s 201ms/step
```

```
fig, ax = plt.subplots(ncols= 4, nrows=4, figsize=(10,10))
for r in range(4):
    for c in range(4):
        ax[r][c].imshow(imgs[(r+1)*(c+1)-1])
```



## ✓ Saving the model

```
generator.save_weights('generator_weights.h5')
```

## ✓ Loading the saved model

```
from tensorflow.keras.models import load_model
loaded_model = build_generator()
```

```
loaded_model.compile()
```

```
loaded_model.load_weights('generator_weights.h5')
```

```
loaded_model_img = loaded_model.predict(tf.random.normal((16, 128, 1)))
fig, ax = plt.subplots(ncols= 4, nrows=4, figsize=(10,10))
for r in range(4):
    for c in range(4):
        ax[r][c].imshow(loaded_model_img[(r+1)*(c+1)-1])
```

1/1 [=====] - 0s 121ms/step

